

Lecture 10 Bayesian DSGE Models

Fei Tan

Department of Economics
Chaifetz School of Business
Saint Louis University

Introduction to Bayesian Statistics

March 19, 2025

Introduction

- ▶ DYNARE has played a large role in fitting DSGE models
- ▶ However, it cannot handle *high-dimensional* DSGE models or models with
 - ▶ Student- t shocks, e.g., Chib & Ramamurthy (2014)
 - ▶ stochastic volatility, e.g., Justiniano & Primiceri (2008)
- ▶ Chib, Shin, & Tan (2021) provide a user-friendly MATLAB toolbox for such models that contains
 - ▶ training sample priors
 - ▶ efficient sampling of parameters by the TaRB-MH algorithm of Chib & Ramamurthy (2010)
 - ▶ fast computation of the marginal likelihood by the method of Chib (1995) and Chib & Jeliazkov (2001)
 - ▶ post-estimation tools, e.g., point and density forecasts

The Road Ahead...

- 1 DSGE Model
- 2 Prior Distribution
- 3 Posterior Sampling
- 4 Marginal Likelihood
- 5 Prediction

Illustration

- Dynamic IS relation

$$\hat{y}_t = \mathbb{E}_t \hat{y}_{t+1} + \hat{g}_t - \mathbb{E}_t \hat{g}_{t+1} - \frac{1}{\tau} (\hat{R}_t - \mathbb{E}_t \hat{\pi}_{t+1} - \mathbb{E}_t \hat{z}_{t+1})$$

- New Keynesian Phillips curve

$$\hat{\pi}_t = \beta \mathbb{E}_t \hat{\pi}_{t+1} + \kappa (\hat{y}_t - \hat{g}_t)$$

- Monetary policy rule

$$\hat{R}_t = \rho_R \hat{R}_{t-1} + (1 - \rho_R) \psi_1 \hat{\pi}_t + (1 - \rho_R) \psi_2 (\hat{y}_t - \hat{g}_t) + \epsilon_{R,t}$$

- Exogenous shocks

$$\hat{g}_t = \rho_G \hat{g}_{t-1} + \epsilon_{G,t}, \quad \hat{z}_t = \rho_Z \hat{z}_{t-1} + \epsilon_{Z,t}$$

- t -innovation with SV: for $s \in \{R, G, Z\}$

$$\epsilon_{s,t} \sim t_\nu(0, e^{h_{s,t}}), \quad h_{s,t} = (1 - \phi_s) \mu_s + \phi_s h_{s,t-1} + \eta_{s,t}, \quad \eta_{s,t} \sim N(0, \omega_s^2)$$

State Space Form

- ▶ We rely on Sims' (2001) method to solve for the solution of the structural model, for *each* value of θ
- ▶ After (log) linearizing around the steady state, the structural model is expressed in canonical form as

$$\Gamma_0(\theta)x_t = \Gamma_1(\theta)x_{t-1} + \Psi\epsilon_t + \Pi\eta_t$$

- ▶ Applying Sims' method, the unique, bounded solution takes the form

$$x_t = G(\theta)x_{t-1} + M(\theta)\epsilon_t$$

where $G(\theta)$ and $M(\theta)$ are non-linear unspecified functions of θ determined by the solve step

State Space Form

- Model completed by the measurement equations

$$\begin{pmatrix} \text{YGR}_t \\ \text{INF}_t \\ \text{INT}_t \end{pmatrix} = \begin{pmatrix} \gamma^{(Q)} \\ \pi^{(A)} \\ \pi^{(A)} + r^{(A)} + 4\gamma^{(Q)} \end{pmatrix} + \begin{pmatrix} \hat{y}_t - \hat{y}_{t-1} + \hat{z}_t \\ 4\hat{\pi}_t \\ 4\hat{R}_t \end{pmatrix}$$

- We show results for an extended version of this model that has 51 parameters, 21 variables, 8 shocks, 8 observables, and 1,494 non-Gaussian and nonlinear latent variables

Example

- ▶ Set MATLAB directory to the 'DSGE-SVt' folder
- ▶ Specify the model, data, and save directories

```
%% Housekeeping
clear
close all
clc

%% User search path & mex files
modpath = ['user' filesep 'ltw17'];
datapath = ['user' filesep 'ltw17' filesep '
    data.txt'];
savepath = ['user' filesep 'ltw17'];
spec = tarb_spec([], 'modpath', modpath, 'datapath
    ', datapath, 'savepath', savepath);

OneFileToMexThemAll
```

The Road Ahead...

- 1 DSGE Model
- 2 Prior Distribution**
- 3 Posterior Sampling
- 4 Marginal Likelihood
- 5 Prediction

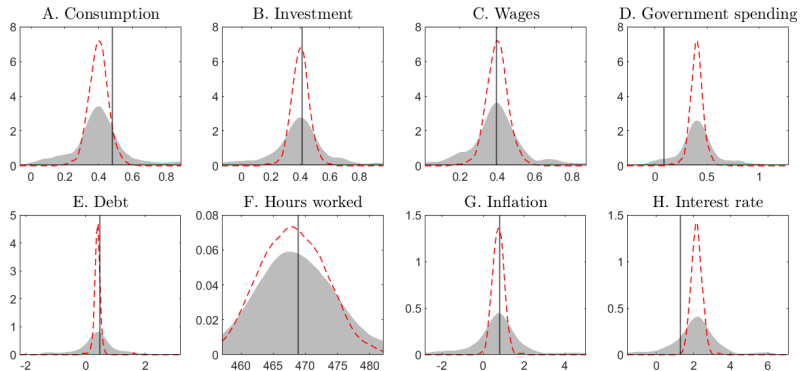
Training Sample Prior

- ▶ In high dimensions, formulating an appropriate prior is difficult due to the complex mapping from θ to $G(\theta)$ and $M(\theta)$
- ▶ Standard choices often produce prior-sample conflict
- ▶ We supply two ways of dealing with this: training sample priors and Student- t family of distributions for location-type parameters
- ▶ A sampling the prior function is available to calculate the implied distribution of the outcomes

```
%% Sampling the prior
npd = 10000;           % number of prior draws
T = 200;               % number of periods
sdof = 5;              % shock degrees of freedom
SamplePrior(npd,T,sdof,savepath)
```

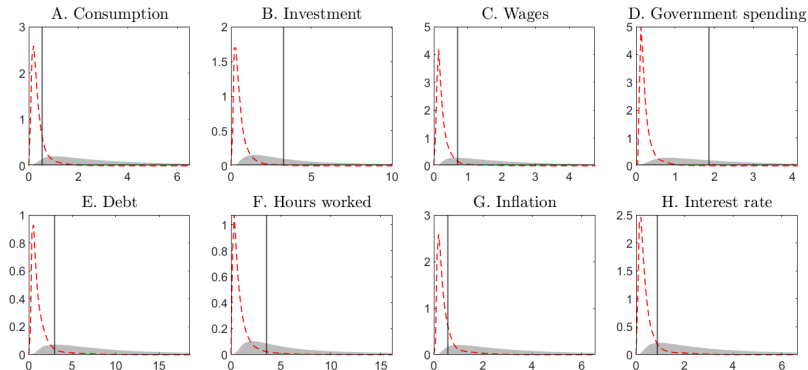
Prior Predictive Distribution

I. Sample mean



Prior Predictive Distribution (Cont'd)

II. Sample standard deviation



Prior Predictive Distribution (Cont'd)

III. Sample first-order autocorrelation



The Road Ahead...

- 1 DSGE Model
- 2 Prior Distribution
- 3 Posterior Sampling**
- 4 Marginal Likelihood
- 5 Prediction

Estimation

- ▶ The centerpiece of the estimation procedure is the TaRB-MH algorithm of Chib & Ramamurthy (2010)
- ▶ Used to draw samples of $(\theta, z_{1:T})$ from

$$\pi(\theta, z_{1:T} | y_{1:T}) \propto f(y_{1:T}, z_{1:T} | \theta) \cdot \pi(\theta) \cdot \mathbf{1}\{\theta \in \Theta_D\}$$

- ▶ TaRB-MH is coded up in DYNARE, but the implementation there is somehow not efficient

TaRB-MH MCMC

- ▶ Two hallmarks of TaRB-MH
 - ▶ randomize the number and components of blocks
 - ▶ tailor the proposal density to the posterior location and curvature, i.e., $q_b(\theta_b|y_{1:T}, z_{1:T}, \theta_{-b}) = t_v(\theta_b|\hat{\theta}_b, \hat{V}_b)$, $b = 1, \dots, B$, using simulated annealing and Chris Sims' `csminwel`
- ▶ The toolbox provides a fast implementation of these steps by
 - ▶ randomly tailoring proposal every few iterations
 - ▶ MEX loop-intensive functions from C/C++ code
 - ▶ `csminwel` to get \hat{V}_b as output
- ▶ Although intensive it all happens seamlessly


Example

- ▶ Set up the TaRB-MH algorithm
- ▶ The estimation results are stored in the MATLAB data file `tarb_full.mat`









```
%% TaRB-MH (full run)
p = 0.7;                % blocking probability
w = 0.5;                % tailoring frequency
spec = tarb_spec(spec, 'prob', p, 'freq', w);

M = 11000;              % number of draws
burn = 1000;            % number of burn-in
tarb_full(M, burn, spec)
```



Prior-Posterior



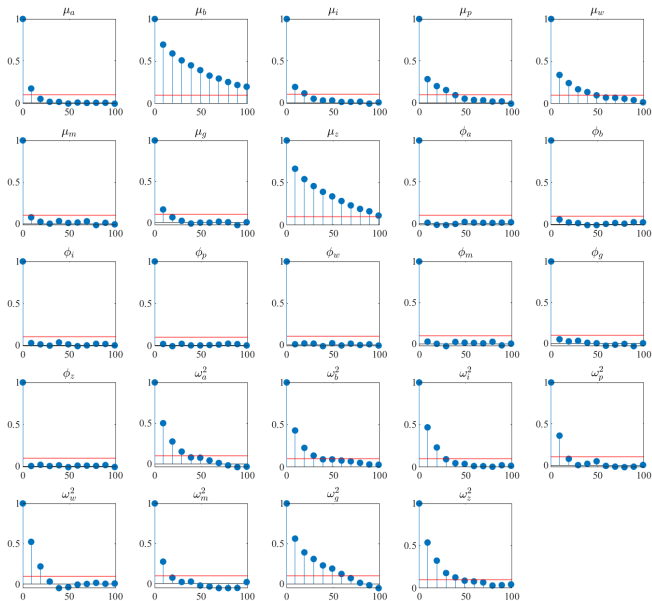
Stochastic Volatility




Autocorrelation



Autocorrelation (Cont'd)



Inefficiency Factor



The Road Ahead...

- 1 DSGE Model
- 2 Prior Distribution
- 3 Posterior Sampling
- 4 Marginal Likelihood**
- 5 Prediction

Model Comparison

- ▶ Marginal likelihood is computed by a fast implementation of the Chib & Jeliazkov (2001) estimator
- ▶ This estimator is based on the identity of Chib (1995)

$$m(y_{1:T}) = \frac{f(y_{1:T}|\theta^*)\pi(\theta^*)}{\pi(\theta^*|y_{1:T})}, \quad \forall \theta^*$$

Model Comparison (Cont'd)

- ▶ With multiple block sampling, the posterior ordinate is estimated from the decomposition

$$\pi(\theta^*|y_{1:T}) = \pi(\theta_1^*|y_{1:T})\pi(\theta_2^*|y_{1:T}, \theta_1^*) \cdots \pi(\theta_B^*|y_{1:T}, \theta_1^*, \dots, \theta_{B-1}^*)$$

using several reduced TaRB-MH runs

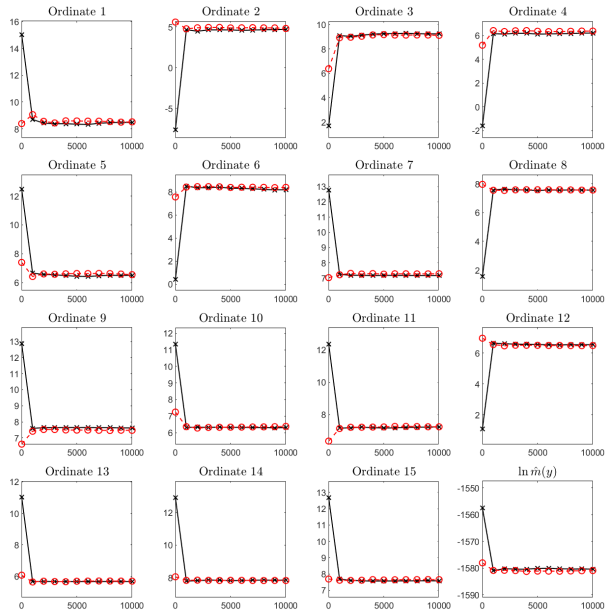
- ▶ Key point: all reduced runs can be parallelized for the cost of only one reduced run, regardless of the number of blocks
- ▶ This speeds up the calculation enormously

Example

- ▶ Set up the TaRB-MH reduced run
- ▶ The estimation results are stored in the MATLAB data file `tarb_reduce.mat`

```
%% TaRB-MH (reduced run)
M = 11000;           % number of draws
burn = 1000;         % number of burn-in
B = 7;               % number of blocks
tarb_reduce(M,burn,B,spec)
```

Marginal Likelihood



Simulation Evidence

Table 1: Chib-Jeliazkov estimator

DGP 1: regime-M with $\nu = 15$			DGP 2: regime-F with $\phi = 0.5$		
ν	M	F	ϕ	M	F
30 (light)	4	0	0.1 (weak)	0	9
15 (fat)	15	0	0.5 (moderate)	0	10
5 (heavy)	1	0	0.9 (strong)	0	1

NOTES: Number of picks for each model specification.

Table 2: Modified harmonic mean estimator

DGP 1: regime-M with $\nu = 15$			DGP 2: regime-F with $\phi = 0.5$		
ν	M	F	ϕ	M	F
30 (light)	0	0	0.1 (weak)	0	0
15 (fat)	0	0	0.5 (moderate)	0	0
5 (heavy)	20	0	0.9 (strong)	0	20

The Road Ahead...

- ① DSGE Model
- ② Prior Distribution
- ③ Posterior Sampling
- ④ Marginal Likelihood
- ⑤ Prediction

Post-Estimation Tools







- ▶ Toolbox includes post-estimation tools, e.g., functions for conducting impulse response, variance decomposition, and prediction
- ▶ For instance, suppose we want the one-quarter-ahead prediction density

$$p(y_{T+1}|y_{1:T}) = \int p(y_{T+1}|y_{1:T}, \theta, z_{T+1}) \cdot \pi(\theta, z_{T+1}|y_{1:T}) d(\theta, z_{1:T})$$

- ▶ This is available through the TaRB-MH specification

```
%% Out-of-sample forecast
head = 50;           % training sample size
tail = 22;           % forecasting sample size
h = 1;               % forecasting horizon
spec = tarb_spec([], 'datrange', [head tail], '
    pred', h);
```

Out-of-Sample Forecast



Outroduction

- ▶ Going beyond DYNARE for Next-Gen DSGE models
- ▶ Efficient and fast estimation via TaRB-MH and parallel computing
- ▶ MATLAB toolbox: publicly available at sites.google.com/a/slu.edu/tanf/research
- ▶ Readily applied. Current application in progress, an open economy DSGE model that is about twice as big as the largest DSGE model estimated to date

References

- ▶ Chib & Ramamurthy (2010), “Tailored Randomized Block MCMC Methods with Application to DSGE Models,” *Journal of Econometrics*
- ▶ Chib, Shin, & Tan (2021), “DSGE-SVt: An Econometric Toolkit for High-Dimensional DSGE Models with SV and t Errors,” *Computational Economics*
- ▶ Sims (2001), “Solving Linear Rational Expectations Models,” *Computational Economics*